

coq-alpha-pearl README

Joshua Grosso
Faculty Mentor: Michael Vanier

September 7, 2021

1 Introduction

This is a Coq formalization of “Functional Pearls: α -conversion is easy” (Altenkirch, 2002). The most important file is `theories/Alpha.v`.

The original paper presents an elegant definition of α -equivalence, along with related algorithms and theorems. By formalizing this material, we present a machine-checkable and -executable version of this work. We hope that this library exposes definitions of α -equivalence, substitution, etc. that are both easy to use computationally and also easy to reason about in proofs.

To our knowledge, all of the main results contained within the paper itself have been formalized. (Proving equivalence to de Bruijn terms is still in progress, but because it is left to the reader in the original paper, we are comfortable sharing this formalization as-is.)

2 Technical Details

2.1 Project Setup

Install the `mathcomp-ssreflect` and `extructures` libraries (see below), and then run `make html` to build the project and its documentation (which annotates important definitions with their counterparts in the original paper).

Detailed installation instructions are as follows: First, install OPAM. Then, at the command line, add the Coq repository to OPAM via `opam repo add coq-released https://coq.inria.fr/opam/released`. Finally, install the project’s dependencies via `opam install coq-mathcomp-ssreflect coq-extructures`. Now, you should be able to build the project and its documentation via `make html`. (At the moment, `theories/Alpha.v` and `theories/Examples.v` may take some time to compile; this is expected.)

2.2 Project Structure

The project is organized on disk as follows:

`theories/Alpha.v` contains the material from the original paper. It is parameterized by the infinite set \mathcal{V} and its corresponding Fresh function.

`theories/AlphaImpl.v` contains two concrete implementations of the formalization, one where \mathcal{V} is the set of finite strings and another where $\mathcal{V} = \mathbb{N}$. We expect the former to be the default choice for practical purposes; we implemented the latter merely as a proof-of-concept.

`theories/Examples.v` contains several examples of \equiv_α in action.

`theories/Util.v` and `theories/Util/*.v` provide a small, custom utility library for the project. Note that `theories/Util/PlfTactics.v` is a copy of `LibTactics.v` from *Programming Language Foundations* (Pierce et al., 2021).

2.3 External Libraries

The project relies on the Mathematical Components library as a supplement to Coq’s built-in standard library. We also rely on `extructures` for extensionally-equal finite sets and maps. Finally, we rely on `LibTactics.v` from *Programming Language Foundations*.

3 Formalization Overview

This formalization contains several deviations from the original paper. Some were for ease of formalization, while others were (as best we can tell) required to prove the desired statements. We tried to be generally faithful to the provided proofs, at least for the more complex theorems. Finally, although we hope it is not the case, any or all of the below changes could have arisen from misinterpretations on our part of the original author’s meaning.

3.1 Typed Lambda Calculus

Our representation of the lambda calculus is untyped, for convenience. However, because the material is agnostic to the presence of types, it should be trivial to adapt it for e.g. the simply-typed λ -calculus.

3.2 (In)finite Sets

We encode the infinite set \mathcal{V} as a type with an order:

Parameter $V : \text{ordType}$.

(Note that defining \mathcal{V} as an `ordType` provides decidable equality as well).

The ordering constraint is not strictly necessary for formalization (and is not present in the original paper), but it allows us to use the `extructures` library’s implementation of extensionally-equal finite sets. This is very convenient for formalization, and we expect that orderings exist for all usual definitions of \mathcal{V} (e.g. the set of all finite strings, or \mathbb{N}).

Of course, if this expectation proves to be incorrect, we can remove this constraint (at the cost of simplicity). For example, we could represent finite sets as linked lists with an appropriate equivalence relation, which we would manually use in lieu of definitional equality.

3.3 Finite Maps

We define partial bijections via the `extructures` library's implementation of extensionally-equal finite maps. We define substitutions similarly; we did not directly encode them as functions because we found that finite maps were more convenient (given that we are avoiding the use of classical logic).

3.4 Inductive Definitions as Boolean Predicates

Where possible, we preferred boolean predicates to inductive propositions (being sure, however, to prove equivalence with the original definitions). For example, consider the definition of $\text{Tm}(X)$ (page 2). A straightforward formalization uses an inductive proposition, like so:

```
(* [x ∈ X] is notation for SSReflect's [x ∈ X]. *)
Section in_Tm.
#[local] Reserved Notation "t '∈' 'Tm' X" (at level 40).

Inductive in_Tm : {fset V} → term →  $\mathbb{P}$  :=
| Tm_variable :  $\forall X x,$ 
  x ∈ X →
  variable x ∈ Tm X
| Tm_application :  $\forall X t u,$ 
  t ∈ Tm X → u ∈ Tm X →
  application t u ∈ Tm X
| Tm_abstraction :  $\forall X t x,$ 
  t ∈ Tm (X ∪ {x}) →
  abstraction x t ∈ Tm X

where "t '∈' 'Tm' X" := (in_Tm X t).
End in_Tm.
```

However, we can alternately define Tm as a boolean predicate (assuming suitable definitions of FV and \subseteq):

Definition $\text{Tm } X \ t : \text{bool} := \text{FV } t \subseteq X$.

Although it is admittedly subjective, we find the latter definition to be semantically clearer. We have proven that these two representations are equivalent:

Lemma $\text{TmP} : \forall X t, \text{reflect } (\text{in_Tm } X \ t) \ (t \in \text{Tm } X)$.

Similarly, we define, use, and prove correctness of algorithmic forms of \equiv_α^R and Tm^{db} . This makes them easy to use in runtime computations.

3.5 Explicitly-Provided Sets

When the original paper introduces a term by writing $t \in \text{Tm}(X)$, we note a slight ambiguity: Given that $\forall X, Y : X \subseteq Y \implies \text{Tm}(X) \subseteq \text{Tm}(Y)$, X is not uniquely specified because any $X \supseteq \text{FV}(t)$ would satisfy the constraint. For ease of formalization (and, to our knowledge, without loss of generality), we have largely replaced references to sets used in this manner with $\text{FV}(t)$. For example, we have formalized Proposition 2.1 as $\forall t : t \equiv_\alpha^{1_{\text{FV}(t)}} t$ instead of $\forall t \in \text{Tm}(X) : t \equiv_\alpha^{1_X} t$. For another example, we defined $t \equiv_\alpha u = t \equiv_\alpha^{1_{\text{FV}(t)}} u$ (noting that $\text{FV}(t) = \text{FV}(u)$ by the assumed α -equivalence of t and u).

We employ a similar strategy for function domains and codomains. Instead of referencing the set X in $f \in X \longrightarrow \text{Tm}(Y)$ (page 4), we instead reference $\text{dom}(f)$. Similarly, instead of referencing Y , we calculate the smallest such set by iterating over the codomain of f (see `codomm_Tm_set`).

Our original approach had been to name these sets explicitly and take them as parameters to the propositions. We were successful in proving the main results of the paper, but the resulting definitions were very unwieldy. Eventually, we revised our work to calculate all such sets where possible, and found that the proof scripts became much shorter as a result.

3.6 Symmetric Updates

Assuming the original paper uses ordered tuples, we have changed the definition of $R(x, y) = (R \setminus \{(x, v), (y, w) \mid v, w \in \mathcal{V}\}) \cup \{(x, y)\}$ to $R(x, y) = (R \setminus \{(x, v), (w, y) \mid v, w \in \mathcal{V}\}) \cup \{(x, y)\}$. We believe this to have been the intention of the original author (given the emphasis on the symmetry of the update procedure).

3.7 Lemma 1.3

We believe that the original proof for Lemma 1.3 (page 3) has a counterexample, arising from the second and third steps of the \iff chain. Specifically, consider the backwards implication portion of that step. Let us separately consider the two cases of our hypothesis $(a = x \wedge z = b) \vee (x \neq a \wedge z \neq b \wedge aR; Sb)$. If $a = x$ and $z = b$, then we can let $c = y$ and our conclusion is satisfied. Thus, we will now assume instead that $x \neq a$, $z \neq b$, and $aR; Sb$. Because R and S are both partial bijections, there must exist a unique c for which aRc and cSb . However, as best we can tell, $c = y$ is not prohibited by the assumptions of the lemma; thus, for the sake of argument, we will assume this equality. In this case, neither case of the conclusion is possible, and so we have a contradiction.

In our formalization, we have replaced $R(x, y); S(y, z) = (R; S)(x, z)$ with $\forall a, b \in \mathcal{V} : aR(x, y); S(y, z)b \implies a(R; S)(x, z)b$. Fortunately, this is still strong enough to prove proposition 2.3, i.e. where lemma 1.3 was originally used.

3.8 Equivalence Classes

Coq does not yet support quotient types (at least not without adding related axioms, which we are hesitant to add). Thus, instead of defining $\text{Tm}(X)/\equiv_\alpha$ at the term-level, we manually proved that the subsequent theorems respect \equiv_α .

3.9 Substitution Extension

In the definition of $\llbracket f \rrbracket$ (page 4), we were unclear whether or not Y is updated upon structural recursion, i.e. whether we should always reference the codomain of the top-level f , or if we should instead use $\text{cod}(f[x, z])$ when recursing in $\llbracket f \rrbracket(\lambda x.t) = \lambda z.\llbracket f[x, z] \rrbracket(t)$. We assumed the latter, and were able to prove the subsequent theorems.

3.10 Substitution Monad Laws

We found it necessary to add an additional condition $x \notin \text{FV}(v)$ to the second monad law for substitution, a la Exercise 2.2 in *Introduction to Lambda Calculus* (Barendregt & Barendsen, 2000).

3.11 Lemma 7

We instead proved $\llbracket f \rrbracket(t) \equiv_\alpha^{f \circ} t$; we assume this was the original author's intention.

3.12 Minor Textual Notes

On page 2, we presume that $X \subseteq_{\text{fin}} \mathcal{V}$ is equivalent (by the infiniteness of \mathcal{V}) to $X \subset_{\text{fin}} \mathcal{V}$.

On page 3, we presume that $R(x, y) \cdots \in (X \cup \{x\}) \times (Y \cup \{y\})$ uses \in to represent \subseteq .

On page 5, we assume that “Hence by ind.hyp. we know $\llbracket f[x, z] \rrbracket(t) \equiv_\alpha^{S(z, z')}$ $\llbracket f[y, z'] \rrbracket(u) \dots$ ” was intended to read “Hence by ind.hyp. we know $\llbracket f[x, z] \rrbracket(t) \equiv_\alpha^{S(z, z')}$ $\llbracket g[y, z'] \rrbracket(u) \dots$ ”.

4 Future Work

Some of our potential next steps:

Proving equivalence to de Bruijn terms is still in progress.

Currently, the definition of `term` is not user-extensible (and is currently rather barebones). We hope to explore ways, if possible, to adapt this formalization to more featureful λ -calculi on demand.

We hope to investigate and improve the runtime performance of our definition of \equiv_α . (Currently, `theories/Examples.v` is rather slow to run.)

Please see the header comment in `theories/Alpha.v` for TODOs relating to the Coq scripts themselves.

5 Acknowledgements

This work was financially supported by the *Soli Deo Gloria* SURF endowment at Caltech.